Aircraft Noise: design summary

What changed on 22 October 2026, why, and what's left · DESE71003 Sensing & IoT coursework

> **Goal.** Two microphones hear an aircraft; the ADS-B receiver identifies it (the aircraft whose sound would arrive at that moment); the Sonos announces it with its route. **Only aircraft that were actually heard are announced** (v1), and later only unusual ones, judged from how often each airline, type and route is seen (v2).

1. Where the system stands (collection goes live tonight)

- **Two mic nodes** (ESP32 + INMP441, mains powered) turn every second of sound into nine numbers and send them, **encrypted**, over ESP-NOW long-range radio to a **bridge** ESP32 on the Pi's USB. No audio leaves a node.
- **The Pi** runs two services: `bridge_link` stores the levels; `adsb_logger` stores every aircraft position within 60 km plus one **pass** row per aircraft within 5 km (the ground truth the mics are checked against). Everything goes into one SQLite database. `analysis/summary.py` produces the time-series figures. The software flow is in Appendix A.
- **Fresh start:** the first 2.2 days of data are archived on the Mac (outside git); the 27 h of raw bridge audio from the first trial were deleted, not archived. Code went from about 60 files to 27, with one README.

2. Recording efficiency: from audio to levels

Recording method	On the air	Stored per mic per day	Privacy
Raw audio, 16 kHz × 16 bit (first trial)	32,000 B/s	2.76 GB	Records speech
Raw audio as 8-bit μ-law (tried, to fit the CH340 USB chip)	16,000 B/s	1.38 GB	Records speech
Now: 1 s levels (LAeq, LAmax, 7 octave bands, flags)	**44 B/s**	**≈ 11.7 MB**	**Speech can't be recovered**

727× less radio data and about 235× less storage per mic than raw audio, with the analysis done where the sound is. Each node filters 16,000 samples a second (A-weighting and seven octave-band filters, about 20 small filter stages per sample), which uses only a few percent of the ESP32. Octave-band levels are what airport noise monitors record, so the data stays useful for telling aircraft from other sounds and for training a model. ADS-B adds about 50 MB a day; the Pi has 49 GB free.

3. Why ESP-NOW rather than the home Wi-Fi, and how it scales

	ESP-NOW long range + bridge (chosen)	Home Wi-Fi straight to the Pi
Reach	Long-range mode: up to ~2× the range of normal Wi-Fi in open air; doesn't depend on the router	Weak at the windows, where the mics sit outside
Independence	Its own link: router restarts, password changes or busy Wi-Fi don't affect it	Depends on the router, DHCP and credentials in the firmware
Security	AES-CCMP encryption; the bridge accepts only the listed nodes; the household network is never touched	WPA2 of the home network; nodes join that network
Modularity	Nodes, relay and Pi are separate parts with one simple packet format	Fewer boards, but nodes tied to the house network
Cost	One extra ESP32 (the bridge) and one USB cable	None

Scalability. Each node sends one 44-byte packet a second, which at the long-range bit rate (¼–½ Mbit/s) takes a few milliseconds: under 1% of the air time, so the radio channel could carry well over a hundred nodes. The tighter limits are: **encrypted peers** on the bridge (7 by default, 17 at most, 20 peers in total), so about 7 nodes as configured and 17 with one setting changed; and the **USB serial link** at 115,200 baud (about 11.5 kB/s, roughly 50 nodes at one ~210-byte JSON line each per second). Beyond that, add a second bridge. On the Pi each node adds about 11.7 MB a day and one database row a second, which SQLite handles easily. Adding a node means one line in `protocol.h`, one in the Pi's mic table, and a reflash.

4. Sound propagation model

An aircraft is treated as a point source. Its sound spreads over a sphere, losing 6 dB per doubling of distance, and the air absorbs a little more with every kilometre:

$$L_{A,max}(r) = L_1 - 20\log_{10}(r / 1\ \text{km}) - \alpha\,(r - 1\ \text{km})$$

r is the slant distance at the aircraft's closest point (from ADS-B), $\mathbf{L_1}$ the level at 1 km and **α** the air absorption in dB/km. Both are fitted per mic, by least squares, to the passes that clearly rose above the background, after subtracting the background's energy (as in ISO 1996). Without that step, distant passes whose "peak" is really background noise flatten the curve. The sound arrives **r / 343 m/s** after the closest moment (3.2 s at 1.1 km), so that's where each pass's level is read. The model is used three ways:

- **Predict and emulate before collecting:** six hours of synthetic data were generated from the model with known values (α = 4.0 dB/km). The analysis recovered α = 4.1–4.4 dB/km, which checked the whole pipeline end to end.
- **Attribution:** the sound delay tells the event detector which aircraft a rise in level belongs to.
- **Insight:** the fitted curve and its scatter show how far away planes are still audible at each mic, and which aircraft are louder than the model predicts (a feature for v2's "special aircraft").

Air absorption rises steeply with frequency (ISO 9613, 10 °C, 70% humidity: about 0.1, 0.4, 1.0, 1.9, 3.7, 9.7 and 33 dB/km in the 63 Hz to 4 kHz octaves), so a pass sounds more muffled the further away it is. The octave bands let that be tested directly: the high bands of a pass should fall faster with distance than the low ones. Simplifications: no Doppler, no ground reflection, and one source level per aircraft type.

5. From mic to announcement: how the pieces fit

Stage	Driven by	What it does
Pass records	ADS-B	Every aircraft within 5 km, heard or not: the ground truth. Lets the report measure what the mics *missed* (detection rate by distance, altitude and window), and labels the training data.
Heard events (to build)	The mics	A rise in level above the background, matched to the pass whose sound would arrive then. v1 uses a rule; v2 a model trained on the octave bands.
Flight info (to build)	Heard events only	Route by callsign from the free, offline **VRS standing data** (updated daily); aircraft type, owner and photo from **adsbdb** (free, cached). No FlightAware or feeding for now: nothing is looked up for flights that weren't heard.
Announcement (to build)	Heard events only	Spoken on the Sonos over whatever is playing (method already proven in `Sonos-Test`). v1: every heard flight, within quiet hours. v2: only the rare ones.

About **460 aircraft a day** pass within 1.5 km, roughly **one every two minutes** in daytime, almost all on the Heathrow approach at about 3,250 ft. That makes v2's rarity filter essential, and gives plenty of training examples within a week. British Airways alone made 454 of 1,044 close passes in the trial, while 25 airlines appeared only once.

6. Data, privacy and security

- **No audio is recorded or stored** anywhere: privacy by design, and the main ethics point. The first trial's audio was deleted.
- **The home location** lives only in `pi/location.json`, which git ignores. Published data will carry distances and bearings from the receiver, never coordinates, because aircraft positions around a receiver reveal where it is.
- **Radio:** AES-CCMP encrypted ESP-NOW. The keys are generated randomly by `flash.sh` and kept in a git-ignored `secrets.h`. The bridge rejects unknown boards and wrong packet versions.
- **The publishing API and dashboard** are deferred until the model is trained, for the deployment part of the report.

7. Robustness built in

- Nodes are mains powered and always listening; they **restart themselves** if the code hangs (watchdog) or the bridge stops acknowledging for 60 s. The bridge sends a heartbeat every minute.
- **Mic-fault flag** when a mic sends nothing but one value. A faulty INMP441 found during testing read pure zeros and crashed the Pi's USB controller; new mics are now tried on a phone charger first.
- Sequence numbers expose lost packets; `uptime` exposes restarts. The Pi logs a health line per mic every 10 minutes and keeps its logs across reboots. Services restart by themselves after a power cut. The ADS-B logger records its own coverage, so an outage can't pass for an empty sky.
- **Calibration** belongs to each mic and lives on the Pi (no reflash). It's done with both nodes in their weatherproof cases, side by side, with pink noise at three fixed levels: an exact A-vs-B offset, an approximate absolute offset from a phone meter, and a check on whether the cases colour any octave band.

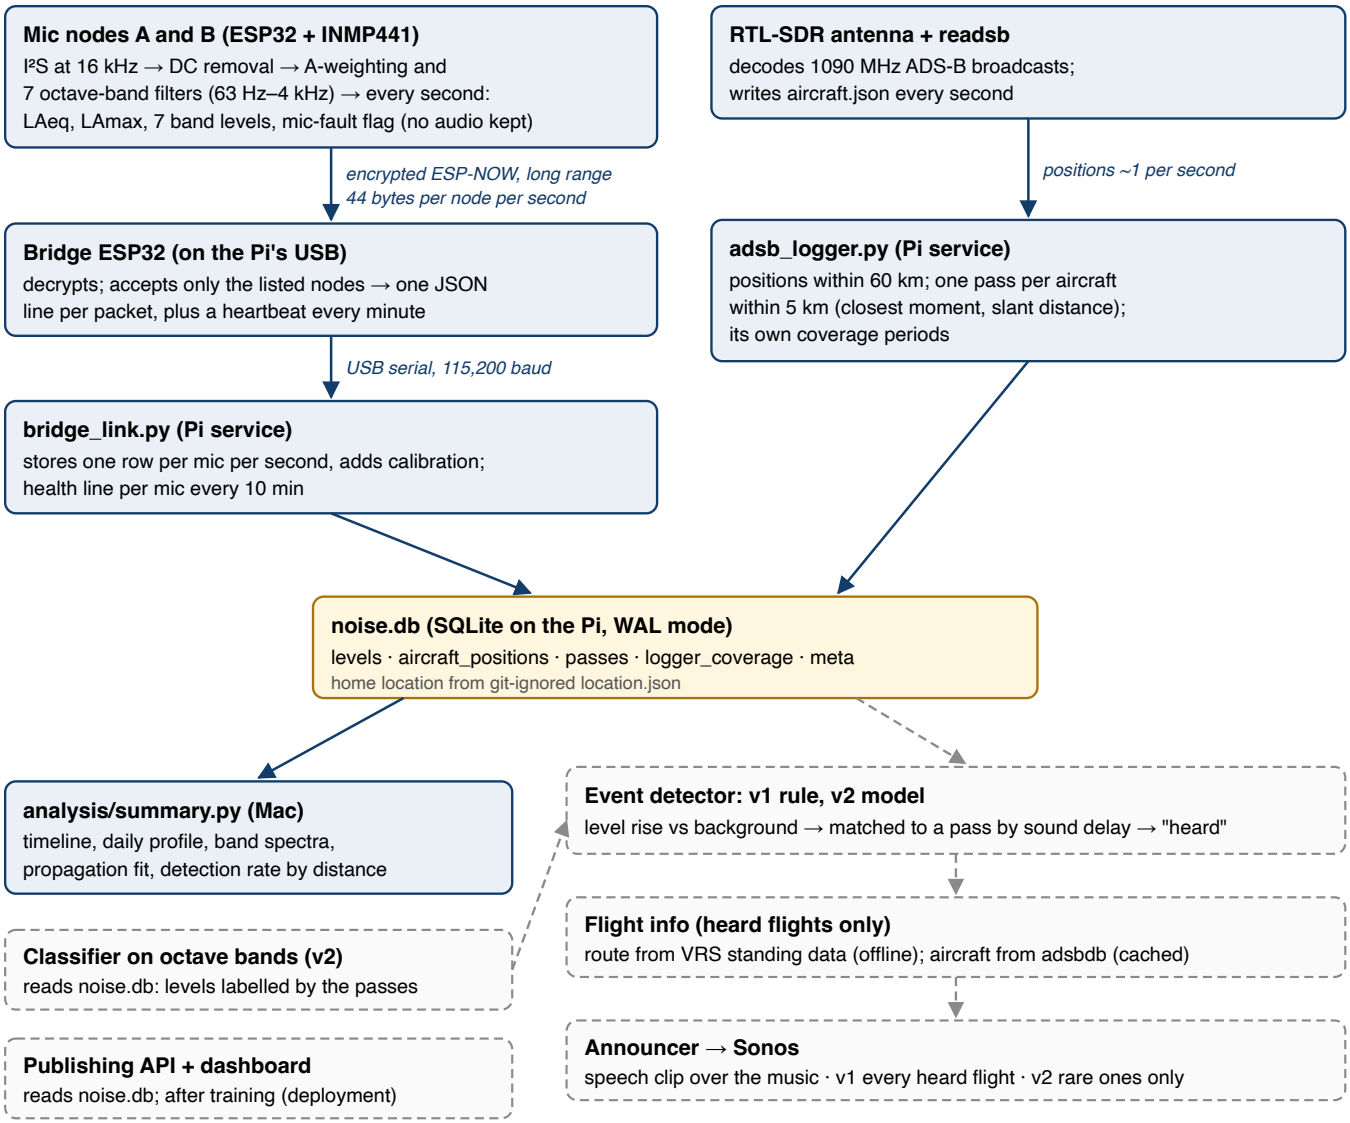
8. v1 then v2

v1 (working system): a rule finds heard aircraft (level rise above the background at the expected arrival time), looks up the route and announces every heard flight. **v2** (learning goal): a classifier trained on the octave-band levels, labelled automatically from the pass records, decides "aircraft or not". Its precision and recall are compared with v1's rule, and only rare aircraft are announced, ranked by how often each airline, type and route appears in the pass history. Both are measured against the same ADS-B ground truth, which gives the report its central analysis.

The full decision log is in Appendix B, and the remaining work, in order, is in Appendix C.

Appendix A. Software workflow

Solid boxes run from tonight. Dashed boxes are still to build (Appendix C).



Appendix B. Decision log

Decision	Why
Store levels, never audio	Privacy, and the module's ethics guidance on sensitive data; 727× less radio data and about 235× less storage. The first trial's audio (27 h) was deleted.
LAeq + LAmax + 7 octave bands, each second	The standard way aircraft noise is monitored; enough detail to tell aircraft from traffic and to train a model on; can't reproduce speech.
Mains-powered, always-on nodes	The nodes are wired to power, so sleep, heartbeats and battery code were removed in favour of robust, continuous data.
Keep the bridge and ESP-NOW (not Wi-Fi)	House Wi-Fi is weak where the mics sit, outside the windows; keeps the sensing network independent of the household network; modular.
ESP-NOW long-range mode	The data is tiny (44 B/s), so the slower bit rate costs nothing and buys range for the untested floor-3 bedroom.
Encrypt ESP-NOW	Security: only the listed nodes are accepted, and keys stay out of git.
Bridge has no mic	Its mic moved to a node. With nothing connected, its data pin floated and reported junk (70–87 dB(A)).
One firmware for both nodes; mics identified by MAC	No per-node builds. Names, locations and calibration live in one table on the Pi, so moving or recalibrating a mic needs no reflash.
Watchdogs, mic-fault flag, health log, persistent logs	Lessons from testing: a dead mic, a crashed USB controller and two power cuts, which were hard to diagnose without logs.
Serial at 115,200 baud	The CH340 USB chip dropped bytes above about 40 kB/s while audio was being sent; levels need under 1 kB/s.
Remove the 50 Hz high-pass; drop the unweighted level	A-weighting and the band filters already reject sub-audible rumble; fewer stages, same information.
A-weighting (dB(A))	Matches what people hear and what a phone meter shows; unweighted levels were 10–15 dB higher in rumbly rooms.
ADS-B antenna stays indoors	The trial showed complete coverage to 60 km in all 12 directions, with positions every ~1 s for nearby aircraft; outdoors would only add far-away range.
Pass records (ADS-B) and heard events (mics)	Passes give the ground truth needed to measure misses and to label training data; mic events drive the lookups and announcements.
Routes from VRS standing data; aircraft from adsbdb	Free, tested (BAW906 → Heathrow–Frankfurt), and VRS works offline. FlightAware's free allowance ($5–10/month) wouldn't cover ~460 passes a day; no feeding for now.
v1 rule, then v2 model	A working system first; the model is the learning goal, judged against the rule on the same ground truth.
v2 announces rare aircraft only	About one pass every two minutes in daytime would make "announce everything" unbearable; rarity is learned from the pass history.
Fresh database; first run archived	A clean dataset with one schema; the 2.2-day trial is kept on the Mac (outside git) for reference.
Location in a git-ignored file	Code and data can be published on GitHub without revealing the home.
Publishing API and dashboard later	They belong to the deployment section, once the model exists.
Simplify: ~60 files → 27, one README	Every one-off test tool, unused option and extra board type removed; easier to reason about and to report.

Appendix C. What's left to do, in order

Tonight: start collecting

1. **Flash node A, then node B.** Plug each into the Pi in turn (or the Mac) and run `esp32/flash.sh node pi` (or `esp32/flash.sh node` on the Mac). The bridge is already flashed. All three boards must use the same encryption keys, which are in `esp32/common/include/secrets.h` on the Mac: don't delete that file.
2. **Switch the Pi to the new services:** `ssh -t sebpi@sebpi sudo bash aircraft-noise/setup/install.sh` (needs your password).
3. **Start:** clear any old-format database the old services created, then remove the pause file (`ssh sebpi@sebpi rm aircraft-noise-data/PAUSED`).
4. **Check after 10 minutes:** the health log should show both mics with few lost packets and no mic faults (`journalctl -u bridge-link`).

This week

5. **Calibrate:** both nodes in their cases, side by side, pink noise at about 50, 65 and 80 dB(A), a minute each, phone meter alongside. Put the offsets in the mic table and redeploy.
6. **Install the nodes** at their windows, then check the radio after a day (signal strength and lost packets). The floor-3 bedroom is the one to watch.
7. **Email Dr Boyle** briefly about sensing in a shared home (no audio recorded), as the brief asks for anything involving people's data.
8. **Keep a note of generative AI use** (Claude Code for code and design discussion), for the required notice.

While the data collects (weeks 2–4)

9. **Event detector (v1):** per mic, find rises in level above the background; match each to the pass whose sound would arrive then; mark the pass as heard. Tune the threshold on the first week's data.
10. **Flight info:** download the VRS route files weekly; look up the route of each heard flight by callsign; aircraft type and owner from adsbdb, cached per aircraft.
11. **Announcer (v1):** reuse the Sonos-Test clip method; generate the spoken sentence; quiet hours and a minimum gap.
12. **Classifier (v2):** build features from the octave bands around each moment, with labels from the passes (aircraft, no aircraft, ambiguous dropped); train and compare with the rule (precision, recall).
13. **Rarity (v2):** score each heard flight by how often its airline, type and route appear; announce only rare ones.

Deployment and report (weeks 5–7; deadline 10 December, 4 pm)

14. **Publishing API and dashboard** on the Pi (read-only JSON and a web page: live levels, recent heard flights).
15. **Analysis** on at least a week of data: run `analysis/summary.py`, plus detection performance, announcement delay, packet loss and uptime, data volume, scalability, security and privacy.
16. **Publish the code and data on GitHub:** export the dataset with distances and bearings only (no coordinates).
17. **Write the report** (cover page and 6 pages, IEEE or ACM two-column) and record the 5-minute video. Include the AI-use notice.

Suggested dates: collecting from 23 Oct · one week of data by 30 Oct · detector and flight info by 7 Nov · announcer by 14 Nov · classifier and rarity by 21 Nov · API, dashboard and analysis by 28 Nov · report and video by 7 Dec.